

In [2]:

```
# Importig Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.preprocessing import StandardScaler

from sklearn import svm
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report

from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split, cross_val_score

# Loading data
train_data = pd.read_csv('SalaryData_Train(1).csv')
test_data = pd.read_csv('SalaryData_Test(1).csv')
```

In [3]:

```
#EDA & Data Preprocessing
train_data.shape
```

Out[3]:

(30161, 14)

In [4]:

```
test_data.shape
```

Out[4]:

(15060, 14)

In [5]:

```
train_data.head()
```

Out[5]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	39	State-gov	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female

In [6]:

```
test_data.head()
```

Out[6]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White	Male

In [7]:

```
# Checking for null values
train_data.isna().sum()
```

Out[7]:

```
age                0
workclass          0
education          0
educationno        0
maritalstatus      0
occupation         0
relationship       0
race               0
sex                0
capitalgain        0
capitalloss        0
hoursperweek       0
native             0
Salary             0
dtype: int64
```

In [8]:

```
test_data.isna().sum()
```

Out[8]:

```
age                0
workclass          0
education          0
educationno        0
maritalstatus      0
occupation         0
relationship       0
race               0
sex                0
capitalgain        0
capitalloss        0
hoursperweek       0
native             0
Salary             0
dtype: int64
```

In [9]:

```
train_data.dtypes
```

Out[9]:

```
age                int64
workclass          object
education          object
educationno        int64
maritalstatus      object
occupation         object
relationship       object
race              object
sex               object
capitalgain        int64
capitalloss        int64
hoursperweek       int64
native            object
Salary            object
dtype: object
```

In [10]:

```
# frequency for categorical fields
category_col = ['workclass', 'education', 'maritalstatus', 'occupation', 'relationship', 'race']
for c in category_col:
    print(c)
    print(train_data[c].value_counts())
    print('\n')
```

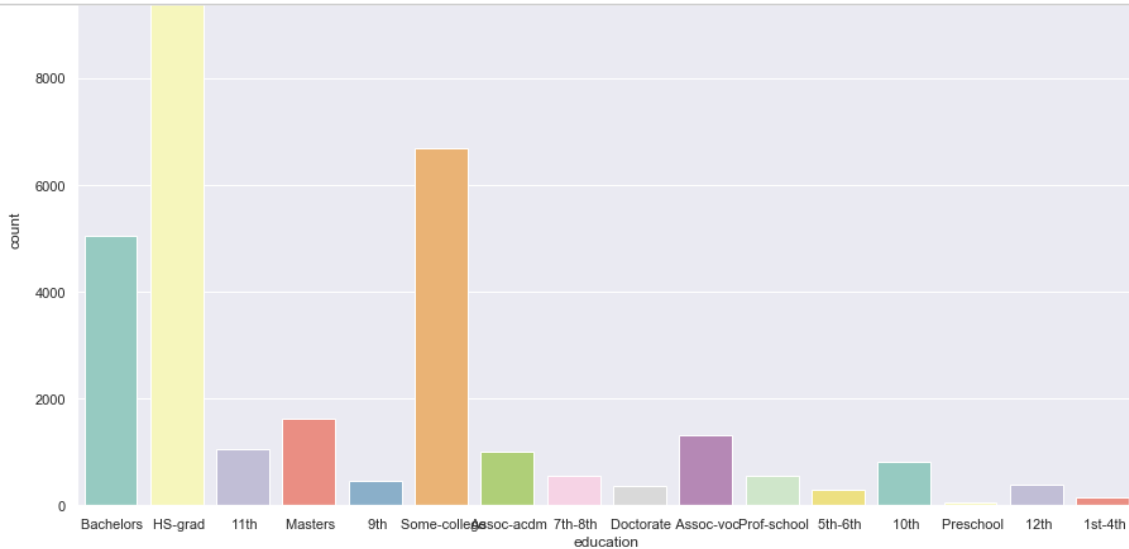
```
1st-4th          151
Preschool        45
Name: education, dtype: int64
```

```
maritalstatus
Married-civ-spouse    14065
Never-married         9725
Divorced              4214
Separated             939
Widowed              827
Married-spouse-absent  370
Married-AF-spouse      21
Name: maritalstatus, dtype: int64
```

```
occupation
Prof-specialty    4038
Craft-repair      4030
Ever-managerial   3997
```

In [11]:

```
# countplot for all categorical columns
import seaborn as sns
sns.set(rc={'figure.figsize':(15,8)})
cat_col = ['workclass', 'education', 'maritalstatus', 'occupation', 'relationship', 'race',
for col in cat_col:
    plt.figure() #this creates a new figure on which your plot will appear
    sns.countplot(x = col, data = train_data, palette = 'Set3');
```



In [12]:

```
# printing unique values from each categorical columns
```

```
print('workclass',train_data.workclass.unique())
print('education',train_data.education.unique())
print('maritalstatus',train_data['maritalstatus'].unique())
print('occupation',train_data.occupation.unique())
print('relationship',train_data.relationship.unique())
print('race',train_data.race.unique())
print('sex',train_data.sex.unique())
print('native',train_data['native'].unique())
print('Salary',train_data.Salary.unique())
```

```
workclass [' State-gov' ' Self-emp-not-inc' ' Private' ' Federal-gov' ' Local-gov'
' Self-emp-inc' ' Without-pay']
education [' Bachelors' ' HS-grad' ' 11th' ' Masters' ' 9th' ' Some-college'
' Assoc-acdm' ' 7th-8th' ' Doctorate' ' Assoc-voc' ' Prof-school'
' 5th-6th' ' 10th' ' Preschool' ' 12th' ' 1st-4th']
maritalstatus [' Never-married' ' Married-civ-spouse' ' Divorced'
' Married-spouse-absent' ' Separated' ' Married-AF-spouse' ' Widowed']
occupation [' Adm-clerical' ' Exec-managerial' ' Handlers-cleaners' ' Prof-specialty'
' Other-service' ' Sales' ' Transport-moving' ' Farming-fishing'
' Machine-op-inspct' ' Tech-support' ' Craft-repair' ' Protective-serv'
' Armed-Forces' ' Priv-house-serv']
relationship [' Not-in-family' ' Husband' ' Wife' ' Own-child' ' Unmarried'
' Other-relative']
race [' White' ' Black' ' Asian-Pac-Islander' ' Amer-Indian-Eskimo' ' Other']
sex [' Male' ' Female']
native [' United-States' ' Cuba' ' Jamaica' ' India' ' Mexico' ' Puerto-Rico'
' Honduras' ' England' ' Canada' ' Germany' ' Iran' ' Philippines'
' Poland' ' Columbia' ' Cambodia' ' Thailand' ' Ecuador' ' Laos'
' Taiwan' ' Haiti' ' Portugal' ' Dominican-Republic' ' El-Salvador'
' France' ' Guatemala' ' Italy' ' China' ' South' ' Japan' ' Yugoslavia'
' Peru' ' Outlying-US(Guam-USVI-etc)' ' Scotland' ' Trinidad&Tobago'
' Greece' ' Nicaragua' ' Vietnam' ' Hong' ' Ireland' ' Hungary']
Salary [' <=50K' ' >50K']
```

In [13]:

```
train_data[['Salary', 'age']].groupby(['Salary'], as_index=False).mean().sort_values(by='age')
```

Out[13]:

	Salary	age
1	>50K	43.959110
0	<=50K	36.608264

In [14]:

```
plt.style.use('seaborn-whitegrid')
x, y, hue = "race", "prop", "sex"
#hue_order = ["Male", "Female"]
plt.figure(figsize=(20,5))
f, axes = plt.subplots(1, 2)
sns.countplot(x=x, hue=hue, data=train_data, ax=axes[0])


prop_df = (train_data[x]
           .groupby(train_data[hue])
           .value_counts(normalize=True)
           .rename(y)
           .reset_index())

sns.barplot(x=x, y=y, hue=hue, data=prop_df, ax=axes[1])
```

Out[14]:

<AxesSubplot:xlabel='race', ylabel='prop'>

<Figure size 1440x360 with 0 Axes>



In [15]:

```
g = sns.jointplot(x = 'age',
                  y = 'hoursperweek',
                  data = train_data,
                  kind = 'hex',
                  cmap= 'hot',
                  size=10)
```

```
#http://stackoverflow.com/questions/33288830/how-to-plot-regression-line-on-hexbins-with-se  
sns.regplot(train_data.age, train_data['hoursperweek'], ax=g.ax_joint, scatter=False, color
```

C:\Users\sowmya sandeep\anaconda3\lib\site-packages\seaborn\axisgrid.py:218
2: UserWarning: The `size` parameter has been renamed to `height`; please update your code.


warnings.warn(msg, UserWarning)

C:\Users\sowmya sandeep\anaconda3\lib\site-packages\seaborn_decorators.py:3
6: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

Out[15]:

<AxesSubplot:xlabel='age', ylabel='hoursperweek'>



In [16]:

```
#Feature encoding
from sklearn.preprocessing import LabelEncoder

train_data = train_data.apply(LabelEncoder().fit_transform)
train_data.head()
```

Out[16]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	22	5	9	12	4	0	1	4	1
1	33	4	9	12	2	3	0	4	1
2	21	2	11	8	0	5	1	4	1
3	36	2	1	6	2	5	0	2	1
4	11	2	9	12	2	9	5	2	0

In [17]:

```
test_data = test_data.apply(LabelEncoder().fit_transform)
test_data.head()
```

Out[17]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	8	2	1	6	4	6	3	2	1
1	21	2	11	8	2	4	0	4	1
2	11	1	7	11	2	10	0	4	1
3	27	2	15	9	2	6	0	2	1
4	17	2	0	5	4	7	1	4	1

In [19]:

```
#Test-Train-Split
drop_elements = ['education', 'native', 'Salary']

X = train_data.drop(drop_elements, axis=1)

y = train_data['Salary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```

In [20]:

```
#Building SVM Model
from sklearn import metrics

svc = SVC()
svc.fit(X_train, y_train)
# make predictions
prediction = svc.predict(X_test)
# summarize the fit of the model
print(metrics.classification_report(y_test, prediction))
print(metrics.confusion_matrix(y_test, prediction))

print("Accuracy:", metrics.accuracy_score(y_test, prediction))
print("Precision:", metrics.precision_score(y_test, prediction))
print("Recall:", metrics.recall_score(y_test, prediction))
```

	precision	recall	f1-score	support
0	0.80	0.99	0.89	7466
1	0.86	0.28	0.42	2488
accuracy			0.81	9954
macro avg	0.83	0.63	0.65	9954
weighted avg	0.82	0.81	0.77	9954

```
[[7355 111]
 [1789 699]]
Accuracy: 0.8091219610206952
Precision: 0.8629629629629629
Recall: 0.2809485530546624
```

In [21]:

```
#Testing it on new test data from SalaryData_Test(1).csv
drop_elements = ['education', 'native', 'Salary']
X_new = test_data.drop(drop_elements, axis=1)

y_new = test_data['Salary']

# make predictions
new_prediction = svc.predict(X_new)
# summarize the fit of the model
print(metrics.classification_report(y_new, new_prediction))
print(metrics.confusion_matrix(y_new, new_prediction))

print("Accuracy:",metrics.accuracy_score(y_new, new_prediction))
print("Precision:",metrics.precision_score(y_new, new_prediction))
print("Recall:",metrics.recall_score(y_new, new_prediction))
```

	precision	recall	f1-score	support
0	0.80	0.99	0.89	11360
1	0.87	0.26	0.40	3700
accuracy			0.81	15060
macro avg	0.84	0.63	0.65	15060
weighted avg	0.82	0.81	0.77	15060

```
[[11216  144]
 [ 2727  973]]
Accuracy: 0.8093625498007968
Precision: 0.8710832587287377
Recall: 0.26297297297297295
```

In [22]:

```
#Building SVM model with Hyper Parameters kernel='rbf',gamma=15, C=1
model = SVC(kernel='rbf',gamma=15, C=1)

model.fit(X_train, y_train)
# make predictions
prediction = model.predict(X_test)
# summarize the fit of the model
print(metrics.classification_report(y_test, prediction))
print(metrics.confusion_matrix(y_test, prediction))

print("Accuracy:",metrics.accuracy_score(y_test, prediction))
print("Precision:",metrics.precision_score(y_test, prediction))
print("Recall:",metrics.recall_score(y_test, prediction))
```

	precision	recall	f1-score	support
0	0.76	0.98	0.86	7466
1	0.56	0.08	0.15	2488
accuracy			0.75	9954
macro avg	0.66	0.53	0.50	9954
weighted avg	0.71	0.75	0.68	9954

```
[[7304 162]
 [2280 208]]
Accuracy: 0.754671488848704
Precision: 0.5621621621621622
Recall: 0.08360128617363344
```

In [23]:

```
#Testing above model on SalaryData_Test(1).csv
# make predictions
new_prediction = model.predict(X_new)
# summarize the fit of the model
print(metrics.classification_report(y_new, new_prediction))
print(metrics.confusion_matrix(y_new, new_prediction))

print("Accuracy:",metrics.accuracy_score(y_new, new_prediction))
print("Precision:",metrics.precision_score(y_new, new_prediction))
print("Recall:",metrics.recall_score(y_new, new_prediction))
```

	precision	recall	f1-score	support
0	0.76	0.98	0.86	11360
1	0.55	0.07	0.13	3700
accuracy			0.76	15060
macro avg	0.66	0.53	0.49	15060
weighted avg	0.71	0.76	0.68	15060

```
[[11147 213]
 [ 3437 263]]
Accuracy: 0.7576361221779548
Precision: 0.5525210084033614
Recall: 0.07108108108108108
```

In [24]:

```
#Building SVM model with Hyper Parameters kernel='linear',gamma=0.22, C=0.1
model_2 = SVC(kernel='linear',gamma=0.22, C=1)

model_2.fit(X_train, y_train)
# make predictions
prediction = model.predict(X_test)
# summarize the fit of the model
print(metrics.classification_report(y_test, prediction))
print(metrics.confusion_matrix(y_test, prediction))

print("Accuracy:",metrics.accuracy_score(y_test, prediction))
print("Precision:",metrics.precision_score(y_test, prediction))
print("Recall:",metrics.recall_score(y_test, prediction))
```

	precision	recall	f1-score	support
0	0.76	0.98	0.86	7466
1	0.56	0.08	0.15	2488
accuracy			0.75	9954
macro avg	0.66	0.53	0.50	9954
weighted avg	0.71	0.75	0.68	9954

```
[[7304 162]
 [2280 208]]
Accuracy: 0.754671488848704
Precision: 0.5621621621621622
Recall: 0.08360128617363344
```

In [25]:

```
#Testing above model on SalaryData_Test(1).csv
# make predictions
new_prediction = model_2.predict(X_new)
# summarize the fit of the model
print(metrics.classification_report(y_new, new_prediction))
print(metrics.confusion_matrix(y_new, new_prediction))

print("Accuracy:",metrics.accuracy_score(y_new, new_prediction))
print("Precision:",metrics.precision_score(y_new, new_prediction))
print("Recall:",metrics.recall_score(y_new, new_prediction))
```

	precision	recall	f1-score	support
0	0.81	0.97	0.88	11360
1	0.77	0.29	0.42	3700
accuracy			0.80	15060
macro avg	0.79	0.63	0.65	15060
weighted avg	0.80	0.80	0.77	15060

```
[[11044 316]
 [ 2639 1061]]
Accuracy: 0.8037848605577689
Precision: 0.7705156136528686
Recall: 0.28675675675675677
```

In []: